Aaron Wang

aaronw5@uic.edu · (425) 623-8843 · aaronw5.github.io · Google Scholar · LinkedIn

Education

University of Illinois Chicago Chicago, IL — August 2022 – Present
Ph.D., Physics

University of Washington Seattle, WA — September 2018 – June 2022
B.S., Comprehensive Physics

Awards

DOE Office of Science Graduate Student Research (SCGSR) Award – funds thesis research at a DOE national lab — 2027

URA Visiting Scholar – Universities Research Association award for research at Fermilab — 2026

Chicagoland Computational Traineeship in High Energy Particle Physics (C2 the P2) – two years of full PhD funding for work at Fermilab — 2023 – 2025

IRIS-HEP Fellowship – scientific computing software for the particle physics community — 2021

Washington Research Foundation Nanofabrication Fellowship – quantum computing research — 2019

Research Experience

Amazon Web Services, Fundamental AI Research New York, NY — May 2026 – Present
Applied Scientist Intern

- Research whether LLM agents can follow a wall-clock time budget on long-horizon tasks and use that time effectively, including time spent in tool calls and code execution that the model cannot measure from its own output
- Built time-budgeted evaluations in Factorio and Kaggle machine-learning competitions (MLE-bench) that measure deadline adherence, how much of the budget agents use, and whether they adapt their strategy to it
- Showed that agents meet deadlines but use only about a quarter of their budget; adding a live clock raised output 77% with no missed deadlines
- Developed an RL method (GRPO) that trains agents to follow the budget and adapt their strategy to its length, e.g. different Factorio strategies for short and long budgets
- Introduced trajectory grafting, a data-augmentation method that splices rollouts across budgets, nearly doubling output under long budgets on held-out tasks
- Modified the model architecture to encode elapsed time, raising on-time task completion from 6% to 50%
- Scaled training to multi-node GPU clusters with hundreds of parallel environments, including a 27B-parameter model trained across AWS regions
- First-author paper at the Resource-Aware Agentic AI Workshop @ NeurIPS 2026

Fermi National Accelerator Laboratory Batavia, IL — June 2023 – Present
Graduate Research Fellow

- Member and author of the CMS Collaboration at the CERN Large Hadron Collider
- Develop efficient transformer architectures for classifying particle jets in Large Hadron Collider data
- Designed a physics-inspired linear transformer that uses 80% less memory and runs faster than a standard transformer while matching state-of-the-art accuracy
- Analyzed attention maps and weight distributions to explain attention sparsity in particle transformers; used the sparsity to improve efficiency
- Trained transformer and mixture-of-experts models (Transformer, SwitchHead, Linformer, DeltaNet) on 100-million-particle datasets with distributed training on GPU nodes at Argonne National Laboratory
- Searched LHC data for rare dark matter interactions; built an automated analysis framework

Amazon Web Services, Fundamental AI Research New York, NY — September 2025 – December 2025
Applied Scientist Intern

- Tested whether evolutionary prompt optimizers (GEPA, EvoPrompt) improve LLM accuracy on math and reasoning tasks, across Qwen3, GPT-OSS and Claude models
- Compared optimized and base prompts with statistical tests on accuracy and pass@k
- Compared optimizer performance on stratified and general datasets
- Analyzed token likelihoods along chain-of-thought trajectories to locate failure points in LLM math reasoning

University of Washington Particle Physics Group Seattle, WA — June 2020 – July 2022
Research Assistant

- Applied state-of-the-art machine learning models to particle jet classification, reaching 99% accuracy
- Ranked input-feature importance for jet classification by pruning features on benchmark datasets
- Built autoencoder and GAN models to detect anomalous events at the LHC trigger

- Quantized recursive and recurrent neural networks for FPGA deployment with hls4ml

Pacific Northwest National Laboratory Richland, WA July 2021 – September 2021
Machine Learning Research Intern
- Built few-shot models in PyTorch to predict the solubility of materials
- Analyzed model errors by molecule type and related them to the underlying chemistry
- Cut the input features by 30% with no loss in accuracy, using Spearman correlation to rank feature importance

University of Washington Quantum Materials Lab Seattle, WA March 2019 – July 2022
Research Assistant
- Grew quantum material crystals by flux, vapor and Bridgman methods, including multiple dopings of MnBiTe and $PdTe_2$
- Measured resistivity versus temperature on prepared samples
- Designed and refined growth recipes for new materials with graduate students

IRIS-HEP New York, NY January 2021 – June 2022
Research Fellow
- Extended the Reproducible Open Benchmarks (ROB) platform to run Jupyter notebooks, using papermill
- Built demos introducing ROB to the particle physics community
- Worked with New York University computer science and physics faculty

Washington Nanofabrication Facility Seattle, WA June 2019 – September 2019
Summer Research Fellow
- Fabricated cantilevers in a clean room, designing etch depths and masking layers and verifying them experimentally
- Used the cantilevers to measure fundamental frequencies and study structural transitions in quantum materials

PROJECTS

Distilling Jet Taggers into Physics Equations github.com/aaronw5/jet-tagger-distillation
- Rewrote the latent representations of two jet taggers, JEDI-linear (FPGA-deployable) and Particle Transformer (ParT), as if-statements on jet physics quantities such as jet and subjet masses
- For JEDI-linear, the distilled formulas reproduce the network's decision on 94% of jets (64 particles per jet) with equal classification accuracy
- The same formulas explain 92% of each neuron's variance in the trained network versus 16% in an untrained copy

Interpretability Explorer for a Trigger-Level Jet Tagger github.com/aaronw5/jedi-linear-explorer
- Wrote every neuron of the quantized JEDI-linear tagger as an exact formula of its particle inputs (p_T, $\Delta\eta$, $\Delta\phi$) and matched each to a jet physics quantity
- Built an interactive explorer with sparse autoencoders, transcoders and decision trees distilled from the network, evaluated on 50,000 held-out jets

Machine Learning for Cryogenic Transistor Modeling github.com/aaronw5/Cryo-Transistor-ML
- Trained neural-network surrogates of transistor current–voltage curves to extract BSIM4 compact-model parameters for SkyWater 130 nm transistors at 77 K, validating every fit in circuit simulation (NGSpice)
- Lowered fitting error relative to the published parameters on all 18 test devices (combined relative RMS error 0.260 to 0.214)
- First-author paper at the Sim2Sci Workshop @ NeurIPS 2026

SELECTED PUBLICATIONS

Full list, including CMS Collaboration papers: Google Scholar. *Equal contribution.

1. **Wang, A.**, Zhao, Z., Xia, A., Sun, C., Gandrakota, A., Ngadiuba, J., Cavanaugh, R., & Duarte, J. (2026). Patch Hierarchical Attention Transformer for Efficient Particle Jet Tagging. ***NeurIPS 2026***. arXiv:2605.21789.
2. **Wang, A.**, Madan, N., Sobal, V., Trager, M., Kleinman, M., Mansimov, E., Xia, W., & Soatto, S. (2026). On the Clock: Towards Punctual and Productive Time-Budgeted AI Agents. ***Resource-Aware Agentic AI Workshop @ NeurIPS 2026***. arXiv:2610.10833.
3. **Wang, A.**, et al. (2026). Surrogate Assisted Compact Model Extraction for Commercial CMOS Process Operating at Cryogenic Temperatures. ***Sim2Sci Workshop @ NeurIPS 2026***.
4. Weitz, J., Demler, D., Hawks, B., **Wang, A.**, Tran, N., & Duarte, J. (2026). SNAC-Pack 2.0: Scaled-Out Surrogate Neural Architecture Codesign. ***AutoML 2026*** (Oral). arXiv:2605.16138.
5. Pendiyala, K., Zia, H., Lee, T., Legge, T., De Leon, A. J., Zhao, Z., **Wang, A.**, Gandrakota, A., Ngadiuba, J., Cavanaugh, R., & Duarte, J. (2026). Conditional Capacity and Routing in Mixture-of-Experts Particle Transformers. *arXiv preprint* arXiv:2610.02701.
6. Legge, T.*, **Wang, A.***, Ortiz, J., Limouzi, V., Zhao, Z., Gandrakota, A., Khoda, E. E., Ngadiuba, J., Duarte, J., & Cavanaugh, R. (2025). Why Is Attention Sparse in Particle Transformer? ***Machine Learning and the Physical Sciences (ML4PS) Workshop @ NeurIPS 2025***. arXiv:2512.00210.
7. **Wang, A.**, Zhao, Z., Katel, S., Sahu, V. G., Khoda, E. E., Gandrakota, A., Ngadiuba, J., Cavanaugh, R., & Duarte, J. (2025). Spatially Aware Linear Transformer (SAL-T) for Particle Jet Tagging. *arXiv preprint* arXiv:2510.23641.

8. **Wang, A.**, Gandrakota, A., Ngadiuba, J., Sahu, V., Bhatnagar, P., Khoda, E. E., & Duarte, J. (2024). Interpreting Transformers for Jet Tagging. ***Machine Learning and the Physical Sciences (ML4PS) Workshop @ NeurIPS 2024***. arXiv:2412.03673.
9. Jiang, Q., Palmstrom, J. C., Singleton, J., et al. (2024). Revealing Fermi surface evolution and Berry curvature in an ideal type-II Weyl semimetal. ***Nature Communications*** 15, 2310. doi:10.1038/s41467-024-46633-w.
10. Khoda, E. E., Rankin, D., Teixeira de Lima, R., Harris, P., Hauck, S., Hsu, S.-C., Kagan, M., Loncar, V., Paikara, C., Rao, R., Summers, S., Vernieri, C., & **Wang, A.** (2023). Ultra-low latency recurrent neural network inference on FPGAs for physics applications with hls4ml. ***Machine Learning: Science and Technology*** 4, 025004. arXiv:2207.00559.

Talks and Teaching

Lecture and tutorial, "Transformers," Parts I and II, Machine Learning for Fundamental Physics (ML4FP) School 2026, Georgia Institute of Technology, Atlanta, GA — June 2026
Talk, "Efficient and Interpretable Transformers for Jet Tagging," Fermilab AI Seminar Series, Batavia, IL — March 2026
Talk, "Spatially Aware Linear Transformer (SAL-T) for Efficient Particle Jet Identification," ML4Jets 2025, California Institute of Technology, Pasadena, CA — August 2025
Talk, "Fast RNN Inference on an FPGA," APS Division of Particles and Fields Meeting (DPF21), virtual — July 2021
Talk, "Implementation of Jupyter Notebooks into the Reproducible Open Benchmarks for Data Analysis Platform (ROB)," APS DPF21, virtual — July 2021
Talk, "Implementation of Jupyter Notebooks into the Reproducible Open Benchmarks for Data Analysis Platform (ROB)," ML4Jets 2021, virtual — July 2021
Talk, "Jupyter Notebook Compatibility with ROB," IRIS-HEP Fellow Presentations, virtual — June 2021

Posters

"Patch Hierarchical Attention Transformer for Efficient Particle Jet Tagging," NeurIPS 2026 — December 2026
"On the Clock: Towards Punctual and Productive Time-Budgeted AI Agents," Resource-Aware Agentic AI Workshop @ NeurIPS 2026 — December 2026
"Surrogate Assisted Compact Model Extraction for Commercial CMOS Process Operating at Cryogenic Temperatures," Sim2Sci Workshop @ NeurIPS 2026 — December 2026
"Why Is Attention Sparse in Particle Transformer?," Machine Learning and the Physical Sciences (ML4PS) Workshop @ NeurIPS 2025 — December 2025
"Efficient Transformer Architectures for Jet Tagging," ACAT 2025, Hamburg, Germany — September 2025
"Efficient and Interpretable Transformers for Particle Physics," Fast Machine Learning for Science Conference 2025, ETH Zurich, Switzerland — September 2025
"Interpreting Transformers for Jet Tagging," Machine Learning and the Physical Sciences (ML4PS) Workshop @ NeurIPS 2024 — December 2024
"Interpreting and Accelerating Transformers for Jet Tagging" (with lightning talk), Fast Machine Learning for Science Conference 2024, Purdue University, West Lafayette, IN — October 2024

Technical Skills

LLM & ML Methods: GRPO and RL post-training, RLHF, supervised fine-tuning, reward design, trajectory-level data augmentation, RoPE and architecture modification, hardware-aware neural architecture search, LLM agents and evaluation
ML Frameworks: PyTorch, TensorFlow, Hugging Face, vLLM, SGLang, Megatron, verl, scikit-learn, NumPy, Pandas
Infrastructure: AWS, multi-node GPU clusters, Docker, Kubernetes, distributed training, ROOT
Programming: Python, Bash, Julia, MATLAB, Mathematica, R, Java, LaTeX
Mathematics: Probability theory, statistics, linear algebra, differential equations, multivariable calculus

Mentorship

Timothy Legge — Transformer Interpretability — Jan 2025 – Sep 2026
Now Ph.D. student in Physics, UC San Diego
Victor Limouzi — Attention-Free Transformer Training — Feb 2025 – May 2026
Incoming M.S. Computer Science, Imperial College London
Jacob Ortiz — Attention-Free Transformer Training — Feb 2025 – May 2026
AI Builders Fellow; incoming M.S. Data Science, UC Berkeley
Trevin Lee — Mixture-of-Experts Particle Transformer Training — Aug 2025 – Present
Haris Zia — Switch Particle Transformer Training — Aug 2025 – Present